

Software Design 2

SDN260S

Multithreading

H. Mataifa

Department of Electronic, Electrical
and Computer Engineering

Outline

- Background: what **threads** are
- **Thread states** and life cycle
- Creating and executing **threads**
- **Thread synchronizaton**
- **Producer/Consumer** relationships

Background: Multithreading

- **Concurrency in computers:**

- Computers normally perform several tasks in **parallel/concurrently**, e.g. compile a program, send a file to a printer, receive email over a network, etc.
- **Concurrency**: performing more than one task at the same time (using a single processor or multiple processors)
- Concurrency is usually achieved by means of **threads**
- **Thread**: *a path of execution within a process* (or, a single sequential flow of control within a program)
- A program/process can have more than one **thread**

- **Multithreading:**

- A **program** with **multiple threads of execution**, where **each thread** has its **own method-call stack** and **program counter**, enabling it to execute concurrently with other **threads** while **sharing** with them **application-wide resources**, e.g. **memory**
- Enables tasks to be distributed across multiple processors (if available), increasing **program responsiveness** and **efficiency** (better use of computing resources)

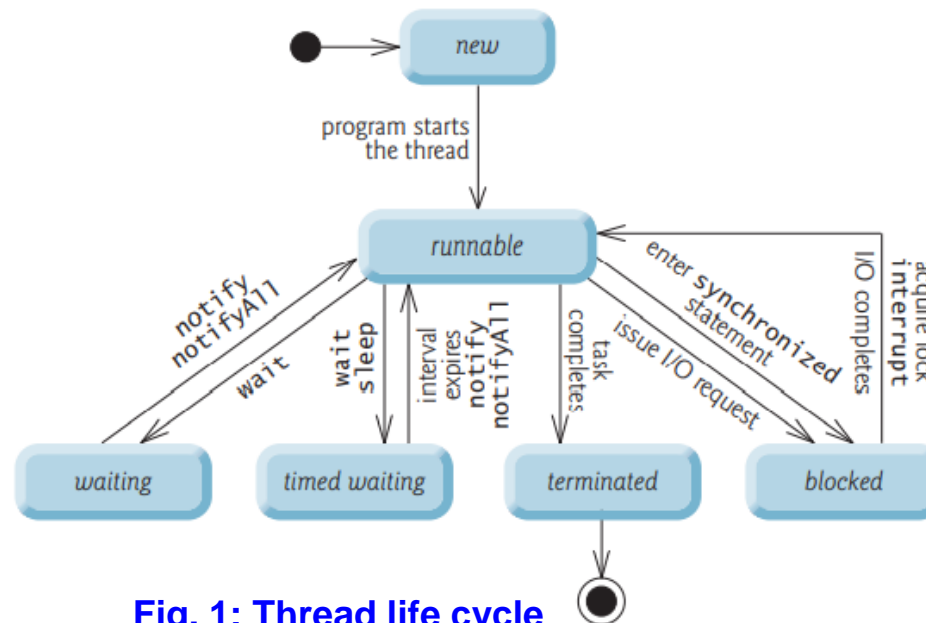
Background: Multithreading

- **Concurrent programming:**
 - A **concurrent program** is essentially a **multithreaded program**; it contains two or more **threads** that **execute concurrently** and work together to perform some task
- **Challenges of multithreaded/concurrent programming:**
 - Has to take care of issues such as **synchronization**, **race conditions**, **deadlock handling**, etc.
- **Support for Concurrent Programming in Java:**
 - To avoid many difficulties and errors associated with concurrent programs, recommended to use Java's **prebuilt concurrency capabilities** (**Concurrency APIs**) to write multithreaded programs
 - “Simple” is better in **multithreaded programming**; use advanced **APIs** (e.g. **Lock, Condition**) only when dictated by application requirements

Thread States: Life Cycle of a Thread

- **Thread states:**

- A new **thread** will be in any one of the states: (1) *runnable*, (2) *waiting*, (3) *timed waiting*, (4) *blocked*, (5) *terminated*
- Various conditions apply that will transition the **thread** from one state to another
- Final state is *terminated*, when the **thread** has successfully completed its designated task, or otherwise terminates due to an error




Waiting for a resource that is being used by another thread

Fig. 1: Thread life cycle

Thread States: Life Cycle of a Thread

- Thread states: **Runnable** state:

- At the operating-system level, thread's **runnable** state typically has two separate states: *ready* and *running*
- *Ready* state can be seen as the entry to **runnable** state (intermediate state before the thread can actually get to run)
- A *ready* thread enters the *running* state (i.e. begins executing) when the operating system assigns it to a processor- also known as *dispatching the thread*
- **Thread scheduling**: the process an operating system uses to determine which thread to **dispatch**
- The operating system uses a **thread priority** (a property of every Java **thread**) in **thread scheduling**; higher-priority **threads** have preference in processor allocation
- Typically, each thread is given a **quantum** or **timeslice** in which to perform a task



Thread is ready to run but it is now running because it is waiting to be dispatched

D

Fig. 2: Thread's runnable state

Thread States: Life Cycle of a Thread

- **Thread scheduling:**

- An operating-system's **thread scheduler** determines which **thread** runs next
- Most operating systems support **timeslicing**, which enables **threads** of equal priority to share a processor, in a **round-robin** fashion, until all **threads** run to completion
- **Preemptive scheduling**: when a higher-priority **thread** enters the **ready state**, the operating system **preempts** the currently running **thread** (essentially reassigns the processor to the higher-priority **thread**)
- **Indefinite postponement/starvation**: could occur when higher-priority **threads** keep postponing the execution of lower-priority **threads**
- **Aging**: employed by the operating system to **prevent starvation** (essentially, a **thread's priority increases with** the amount of **time** it's been waiting)
- **Thread scheduling is platform-dependent**: the behaviour of a multithreaded program could vary across different Java implementations
- Java's **Concurrency APIs** hide much of the issues associated with **Thread Management**

Creating and Executing Threads with **Executor Framework**

- **Interface Runnable:**

- A class implements the **Runnable** interface to be able to perform a task that can execute concurrently with other tasks
- Interface **Runnable** declares the single method **run**, which contains the code that defines the task that a **Runnable object** should perform
- When a **thread** executing a **Runnable object** is created and started, the **thread** calls the **Runnable object's run** method, which executes the task in the new **thread**

- **Executor:**

- Responsible for creating and managing a group of **threads** (**thread pool**)
- **Executor** method **execute** accepts a **Runnable object** as argument, which **Executor** assigns to one of the available **threads** in a **thread pool**
- **Executor framework** provides a means for **efficient thread management**; recommended over custom **thread** creation and management

Creating and Executing Threads with **Executor Framework**

- **Interface ExecutorService:**

- Extends **Executor** interface, declares various methods for managing the life cycle of an **Executor**
- Uses Class **Executors** method `newCachedThreadPool` to create a **thread pool**

- **Class PrintTask:**

- Implements **Runnable** to enable concurrent execution of **PrintTasks**
- **Thread** static method `sleep` places a **thread** in the *timed waiting* state for a specified amount of time
- The program has **main** thread (created by **JVM**) and other **threads** created in main, which execute the **PrintTasks**
- All **threads** (main as well as those created within main) should terminate for the whole program to terminate
- **ExecutorService** method `shutdown` signals termination of the **thread pool**; tasks already submitted are executed to completion, but no new tasks are accepted

Class PrintTask

```
1 // Fig. 26.3: PrintTask.java
2 // PrintTask class sleeps for a random time from 0 to 5 seconds
3 import java.util.Random;
4
5 public class PrintTask implements Runnable
6 {
7     private final int sleepTime; // random sleep time for thread
8     private final String taskName; // name of task
9     private final static Random generator = new Random();
10
11     // constructor
12     public PrintTask( String name )
13     {
14         taskName = name; // set task name
15
16         // pick random sleep time between 0 and 5 seconds
17         sleepTime = generator.nextInt( 5000 ); // milliseconds
18     } // end PrintTask constructor
19
20     // method run contains the code that a thread will execute
21     public void run()
22     {
23         try // put thread to sleep for sleepTime amount of time
24         {
25             System.out.printf( "%s going to sleep for %d milliseconds.\n",
26                               taskName, sleepTime );
27             Thread.sleep( sleepTime ); // put thread to sleep
28         } // end try
29         catch ( InterruptedException exception )
30         {
31             System.out.printf( "%s %s\n", taskName,
32                               "terminated prematurely due to interruption" );
33         } // end catch
34
35         // print task name
36         System.out.printf( "%s done sleeping\n", taskName );
37     } // end method run
38 } // end class PrintTask
```

Class TaskExecutor

```
1 // Fig. 26.4: TaskExecutor.java
2 // Using an ExecutorService to execute Runnable's.
3 import java.util.concurrent.Executors;
4 import java.util.concurrent.ExecutorService;
5
6 public class TaskExecutor
7 {
8     public static void main( String[] args )
9     {
10         // create and name each runnable
11         PrintTask task1 = new PrintTask( "task1" );
12         PrintTask task2 = new PrintTask( "task2" );
13         PrintTask task3 = new PrintTask( "task3" );
14
15         System.out.println( "Starting Executor" );
16
17         // create ExecutorService to manage threads
18         ExecutorService threadExecutor = Executors.newCachedThreadPool();
19
20         // start threads and place in runnable state
21         threadExecutor.execute( task1 ); // start task1
22         threadExecutor.execute( task2 ); // start task2
23         threadExecutor.execute( task3 ); // start task3
24
25         // shut down worker threads when their tasks complete
26         threadExecutor.shutdown();
27
28         System.out.println( "Tasks started, main ends.\n" );
29     } // end main
30 } // end class TaskExecutor
```

Threads are created here

when we do this we are freeing the memory

Starting Executor
Tasks started, main ends

task1 going to sleep for 4806 milliseconds
task2 going to sleep for 2513 milliseconds
task3 going to sleep for 1132 milliseconds
task3 done sleeping
task2 done sleeping
task1 done sleeping

Starting Executor
task1 going to sleep for 3161 milliseconds.
task3 going to sleep for 532 milliseconds.
task2 going to sleep for 3440 milliseconds.
Tasks started, main ends.

task3 done sleeping
task1 done sleeping
task2 done sleeping

Thread Synchronization

- **Thread synchronization:**
 - Coordinated access to shared data by multiple concurrent threads
 - Ensures that only one thread has exclusive access to a shared object that can be modified by all threads
 - Prevents indeterminate results that may occur if access to shared object is not controlled
 - When one thread has exclusive access to the shared object, other threads have to wait before getting access to the object
 - **Mutual exclusion:** ensures that each thread accessing a shared object excludes all others from doing so simultaneously
- **Synchronization via Monitors:**
 - Every object has a monitor and a monitor lock (or intrinsic lock) which can be held by only one thread at a time; can be used to gain exclusive access to the object
 - To modify the object, a thread needs to acquire the lock; other threads attempting to modify the object are blocked until the lock is released
 - A synchronized statement is used to indicate that lock acquisition is needed to access the object; the object is said to be guarded by a monitor lock

Thread Synchronization

- **Synchronized statement:**

```
synchronized ( object )  
{  
    statements  
} // end synchronized statement
```

- **object** is the object whose **monitor lock** needs to be acquired in order to modify it; it is normally **this** if it's the object in which the **synchronized statement** appears
 - The **monitor lock** is automatically released once the **synchronized statement** is exited
 - Java also allows **synchronized methods** (a **monitor lock** must be acquired for the method to be invoked on an object)
- **Unsynchronized data sharing (Programs 26.5-26.7):**
 - **SimpleArray** object (Prog. 26.5) is shared among multiple **threads**, which can all place elements into the array; **Thread** method **sleep** used to highlight issues associated with **unsynchronized access** to the shared data
 - Class **ArrayWriter** (prog. 26.6) implements interface **Runnable** to define a task for inserting elements into **SimpleArray**; task completes after placing three consecutive values into the array
 - Class **SharedArrayTest** (prog. 26.7) executes two **ArrayWriter** tasks that add values to a single **SimpleArray** object
 - **ExecutorService** method **awaitTermination** returns control to its caller either when all tasks complete or when the specified timeout elapses

Unsynchronized Data Sharing: SimpleArray

```
1 // Fig. 26.5: SimpleArray.java
2 // Class that manages an integer array to be shared by multiple threads.
3 import java.util.Arrays;
4 import java.util.Random;
5
6 public class SimpleArray // CAUTION: NOT THREAD SAFE!
7 {
8     private final int[] array; // the shared integer array
9     private int writeIndex = 0; // index of next element to be written
10    private final static Random generator = new Random();
11
12    // construct a SimpleArray of a given size
13    public SimpleArray( int size )
14    {
15        array = new int[ size ];
16    } // end constructor
17
18    // add a value to the shared array
19    public void add( int value )
20    {
21        int position = writeIndex; // store the write index
22
23        try
24        {
25            // put thread to sleep for 0-499 milliseconds
26            Thread.sleep( generator.nextInt( 500 ) );
27        } // end try
28        catch ( InterruptedException ex )
29        {
30            ex.printStackTrace();
31        } // end catch
32
33        // put value in the appropriate element
34        array[ position ] = value;
35        System.out.printf( "%s wrote %2d to element %d.\n",
36                          Thread.currentThread().getName(), value, position );
37
38        ++writeIndex; // increment index of element to be written next
39        System.out.printf( "Next write index: %d\n", writeIndex );
40    } // end method add
41
42    // used for outputting the contents of the shared integer array
43    public String toString()
44    {
45        return "\nContents of SimpleArray:\n" + Arrays.toString( array );
46    } // end method toString
47 } // end class SimpleArray
```

Unsynchronized Data Sharing: **ArrayWriter**

```
1 // Fig. 26.6: ArrayWriter.java
2 // Adds integers to an array shared with other Runnables
3 import java.lang.Runnable;
4
5 public class ArrayWriter implements Runnable
6 {
7     private final SimpleArray sharedSimpleArray;
8     private final int startValue;
9
10    public ArrayWriter( int value, SimpleArray array )
11    {
12        startValue = value;
13        sharedSimpleArray = array;
14    } // end constructor
15
16    public void run()
17    {
18        for ( int i = startValue; i < startValue + 3; i++ )
19        {
20            sharedSimpleArray.add( i ); // add an element to the shared array
21        } // end for
22    } // end method run
23 } // end class ArrayWriter
```

This is to show that there are multiple threads writing and based on the sequence that will determine the output

Unsynchronized Data Sharing: SharedArrayTest

```
1 // Fig 26.7: SharedArrayTest.java
2 // Executes two Runnables to add elements to a shared SimpleArray.
3 import java.util.concurrent.Executors;
4 import java.util.concurrent.ExecutorService;
5 import java.util.concurrent.TimeUnit;
6
7 public class SharedArrayTest
8 {
9     public static void main( String[] arg )
10    {
11        // construct the shared object
12        SimpleArray sharedSimpleArray = new SimpleArray( 6 );
13
14        // create two tasks to write to the shared SimpleArray
15        ArrayWriter writer1 = new ArrayWriter( 1, sharedSimpleArray );
16        ArrayWriter writer2 = new ArrayWriter( 11, sharedSimpleArray );
17
18        // execute the tasks with an ExecutorService
19        ExecutorService executor = Executors.newCachedThreadPool();
20        executor.execute( writer1 );
21        executor.execute( writer2 );
22
23        executor.shutdown();
24
25        try
26        {
27            // wait 1 minute for both writers to finish executing
28            boolean tasksEnded = executor.awaitTermination(
29                1, TimeUnit.MINUTES );
30
31            if ( tasksEnded )
32                System.out.println( sharedSimpleArray ); // print contents
33            else
34                System.out.println(
35                    "Timed out while waiting for tasks to finish." );
36        } // end try
37        catch ( InterruptedException ex )
38        {
39            System.out.println(
40                "Interrupted while waiting for tasks to finish." );
41        } // end catch
42    } // end main
43 } // end class SharedArrayTest
```

By default we do not specify how many threads we will use we can only specify how many tasks

Dispatch the thread and pass it to the cpu

We use a try block to check if the thread is available

Catch any thread that is not available

Unsynchronized Data Sharing: **SharedArrayTest**

```
pool-1-thread-1 wrote 1 to element 0.  
Next write index: 1  
pool-1-thread-1 wrote 2 to element 1.  
Next write index: 2  
pool-1-thread-1 wrote 3 to element 2.  
Next write index: 3  
pool-1-thread-2 wrote 11 to element 0.  
Next write index: 4  
pool-1-thread-2 wrote 12 to element 4.  
Next write index: 5  
pool-1-thread-2 wrote 13 to element 5.  
Next write index: 6  
  
Contents of SimpleArray:  
[11, 2, 3, 0, 12, 13]
```

First pool-1-thread-1 wrote the value 1 to element 0. Later pool-1-thread-2 wrote the value 11 to element 0, thus *overwriting* the previously stored value.

Note:

- One of the challenges of **multithreaded programming** is the spotting of errors – they may occur so infrequently that a ‘broken’ program does not produce incorrect results during testing, creating the illusion that the program is correct

Synchronized Data Sharing: Making Operations Atomic

- **Thread safety:**

- An object is **thread-safe** if it can be accessed by multiple **threads** concurrently without any errors occurring; **SimpleArray** is **not thread-safe**
- **SimpleArray** is **not thread-safe** because it allows any number of **threads** to read and modify shared data concurrently, which tends to cause errors
- To make **SimpleArray** **thread-safe**, method **add** is turned into a **synchronized method**, then at any one time, only one **thread** can modify **SimpleArray**
- **Atomic operation**: a set of operations that cannot be divided into smaller sub-operations; they are placed in a **synchronized statement** to make them **atomic**
- In a **multithreaded program**, all accesses to **mutable** data shared by multiple **threads** are placed in **synchronized statements** or **synchronized methods**

Class SimpleArray with Synchronization

What ever thread
is going to write
or read thread
they will need to
use the add method
which is synchronized

```
1 // Fig. 26.8: SimpleArray.java
2 // Class that manages an integer array to be shared by multiple
3 // threads with synchronization.
4 import java.util.Arrays;
5 import java.util.Random;
6
7 public class SimpleArray
8 {
9     private final int[] array; // the shared integer array
10    private int writeIndex = 0; // index of next element to be written
11    private final static Random generator = new Random();
12
13    // construct a SimpleArray of a given size
14    public SimpleArray( int size )
15    {
16        array = new int[ size ];
17    } // end constructor
18
19    // add a value to the shared array
20    public synchronized void add( int value )
21    {
22        int position = writeIndex; // store the write index
23
24        try
25        {
26            // put thread to sleep for 0-499 milliseconds
27            Thread.sleep( generator.nextInt( 500 ) );
28        } // end try
29        catch ( InterruptedException ex )
30        {
31            ex.printStackTrace();
32        } // end catch
33
34        // put value in the appropriate element
35        array[ position ] = value;
36        System.out.printf( "%s wrote %2d to element %d.\n",
37            Thread.currentThread().getName(), value, position );
38
39        ++writeIndex; // increment index of element to be written next
40        System.out.printf( "Next write index: %d\n", writeIndex );
41    } // end method add
42
43    // used for outputting the contents of the shared integer array
44    public String toString()
45    {
46        return "\nContents of SimpleArray:\n" + Arrays.toString( array );
47    } // end method toString
48 } // end class SimpleArray
```

the synchronised keyword is used to obtain a lock
to the buffer

This keyword is used to achieve
multi exclusion

Class SimpleArray with Synchronization

```
pool-1-thread-1 wrote 1 to element 0.  
Next write index: 1  
pool-1-thread-2 wrote 11 to element 1.  
Next write index: 2  
pool-1-thread-2 wrote 12 to element 2.  
Next write index: 3  
pool-1-thread-2 wrote 13 to element 3.  
Next write index: 4  
pool-1-thread-1 wrote 2 to element 4.  
Next write index: 5  
pool-1-thread-1 wrote 3 to element 5.  
Next write index: 6
```

```
Contents of SimpleArray:  
1 11 12 13 2 3
```

Notes:

- Avoid performing **I/O** operations (or lengthy calculations not requiring synchronization) in **synchronized blocks**; it's important to minimize amount of time that an object is “locked”
- Avoid calling **Thread** method “**sleep**” while holding a lock
- Keep duration of **synchronized statements** as short as possible while maintaining the needed synchronization; this minimizes wait time for blocked **threads**
- Declare data fields that are not expected to change as **final** (can be primitives or object references)

Producer/Consumer Relationship

- **Producer/consumer relationship:**
 - In a producer/consumer relationship, **Producer thread** generates data and stores it in a **shared object**, **Consumer thread** reads data from the **shared object**
 - **Shared object** is called a **buffer**
 - **Producer** writes to **buffer**
 - **Consumer** reads from **buffer**
 - **Producer/consumer relationship** requires **synchronization** to ensure values are produced and consumed properly
 - Operations on **buffer** data shared by **producer/consumer** are **state-dependent**; **buffer** has to be in **correct state** for each operation
 - **Producer** may write to **buffer** if in **not-full state**; **Consumer** may read from **buffer** if in **not-empty state**
 - **Programs 26.9 – 26.13**: details an example of a **Producer/Consumer relationship without synchronization**

Producer/Consumer Relationship: Buffer Interface

```
1  // Fig. 26.9: Buffer.java
2  // Buffer interface specifies methods called by Producer and Consumer.
3  public interface Buffer
4  {
5      // place int value into Buffer
6      public void set( int value ) throws InterruptedException;
7
8      // return int value from Buffer
9      public int get() throws InterruptedException;
10 } // end interface Buffer
```

Producer/Consumer Relationship: Producer Thread

```
1 // Fig. 26.10: Producer.java
2 // Producer with a run method that inserts the values 1 to 10 in buffer.
3 import java.util.Random;
4
5 public class Producer implements Runnable
6 {
7     private final static Random generator = new Random();
8     private final Buffer sharedLocation; // reference to shared object
9
10    // constructor
11    public Producer( Buffer shared )
12    {
13        sharedLocation = shared;
14    } // end Producer constructor
15
16    // store values from 1 to 10 in sharedLocation
17    public void run()
18    {
19        int sum = 0;
20
21        for ( int count = 1; count <= 10; count++ )
22        {
23            try // sleep 0 to 3 seconds, then place value in Buffer
24            {
25                Thread.sleep( generator.nextInt( 3000 ) ); // random sleep
26                sharedLocation.set( count ); // set value in buffer
27                sum += count; // increment sum of values
28                System.out.printf( "\t%2d\n", sum );
29            } // end try
30            // if lines 25 or 26 get interrupted, print stack trace
31            catch ( InterruptedException exception )
32            {
33                exception.printStackTrace();
34            } // end catch
35        } // end for
36
37        System.out.println(
38            "Producer done producing\nTerminating Producer" );
39    } // end method run
40 } // end class Producer
```

Producer/Consumer Relationship: Consumer Thread

```
1 // Fig. 26.11: Consumer.java
2 // Consumer with a run method that loops, reading 10 values from buffer.
3 import java.util.Random;
4
5 public class Consumer implements Runnable
6 {
7     private final static Random generator = new Random();
8     private final Buffer sharedLocation; // reference to shared object
9
10    // constructor
11    public Consumer( Buffer shared )
12    {
13        sharedLocation = shared;
14    } // end Consumer constructor
15
16    // read sharedLocation's value 10 times and sum the values
17    public void run()
18    {
19        int sum = 0;
20
21        for ( int count = 1; count <= 10; count++ )
22        {
23            // sleep 0 to 3 seconds, read value from buffer and add to sum
24            try
25            {
26                Thread.sleep( generator.nextInt( 3000 ) );
27                sum += sharedLocation.get();
28                System.out.printf( "\t\t\t%d\n", sum );
29            } // end try
30            // if lines 26 or 27 get interrupted, print stack trace
31            catch ( InterruptedException exception )
32            {
33                exception.printStackTrace();
34            } // end catch
35        } // end for
36
37        System.out.printf( "\n%s %d\n%s\n",
38            "Consumer read values totaling", sum, "Terminating Consumer" );
39    } // end method run
40 } // end class Consumer
```


Producer/Consumer Relationship: Unsynchronized Buffer

```
1 // Fig. 26.12: UnsynchronizedBuffer.java
2 // UnsynchronizedBuffer maintains the shared integer that is accessed by
3 // a producer thread and a consumer thread via methods set and get.
4 public class UnsynchronizedBuffer implements Buffer
5 {
6     private int buffer = -1; // shared by producer and consumer threads
7
8     // place value into buffer
9     public void set( int value ) throws InterruptedException
10    {
11        System.out.printf( "Producer writes\t%2d", value );
12        buffer = value;
13    } // end method set
14
15    // return value from buffer
16    public int get() throws InterruptedException
17    {
18        System.out.printf( "Consumer reads\t%2d", buffer );
19        return buffer;
20    } // end method get
21 } // end class UnsynchronizedBuffer
```

Producer/Consumer Relationship: SharedBufferTest

```
1 // Fig. 26.13: SharedBufferTest.java
2 // Application with two threads manipulating an unsynchronized buffer.
3 import java.util.concurrent.ExecutorService;
4 import java.util.concurrent.Executors;
5
6 public class SharedBufferTest
7 {
8     public static void main( String[] args )
9     {
10         // create new thread pool with two threads
11         ExecutorService application = Executors.newCachedThreadPool();
12
13         // create UnsynchronizedBuffer to store ints
14         Buffer sharedLocation = new UnsynchronizedBuffer();
15
16         System.out.println(
17             "Action\t\tValue\tSum of Produced\tSum of Consumed" );
18         System.out.println(
19             "-----\t\t\t-----\t\t\t-----\n" );
20
21         // execute the Producer and Consumer, giving each of them access
22         // to sharedLocation
23         application.execute( new Producer( sharedLocation ) );
24         application.execute( new Consumer( sharedLocation ) );
25
26         application.shutdown(); // terminate application when tasks complete
27     } // end main
28 } // end class SharedBufferTest
```

Producer/Consumer Relationship: SharedBufferTest

Action	Value	Sum of Produced	Sum of Consumed	
-----	-----	-----	-----	
Producer writes	1	1		
Producer writes	2	3		—— 1 is lost
Producer writes	3	6		—— 2 is lost
Consumer reads	3		3	
Producer writes	4	10		
Consumer reads	4		7	
Producer writes	5	15		
Producer writes	6	21		—— 5 is lost
Producer writes	7	28		—— 6 is lost
Consumer reads	7		14	
Consumer reads	7		21	—— 7 read again
Producer writes	8	36		
Consumer reads	8		29	
Consumer reads	8		37	—— 8 read again
Producer writes	9	45		
Producer writes	10	55		—— 9 is lost
Producer done producing				
Terminating Producer				
Consumer reads	10		47	
Consumer reads	10		57	—— 10 read again
Consumer reads	10		67	—— 10 read again
Consumer reads	10		77	—— 10 read again
Consumer read values totaling 77				
Terminating Consumer				

Producer/Consumer Relationship: **ArrayBlockingQueue**

- **Class ArrayBlockingQueue:**

- A fully implemented, **thread-safe buffer** class that implements interface **BlockingQueue**
 - Declares methods **put** and **take**, the blocking equivalents of **Queue** methods **offer** and **poll** respectively
- Method **put** places elements at the end of the **BlockingQueue**, waiting if the queue is full
- Method **take** removes an element from the head of the **BlockingQueue**, waiting if the queue is empty

Producer/Consumer Relationship: BlockingBuffer

```
1  // Fig. 26.14: BlockingBuffer.java
2  // Creating a synchronized buffer using an ArrayBlockingQueue.
3  import java.util.concurrent.ArrayBlockingQueue;
4
5  public class BlockingBuffer implements Buffer
6  {
7      private final ArrayBlockingQueue<Integer> buffer; // shared buffer
8
9      public BlockingBuffer()
10     {
11         buffer = new ArrayBlockingQueue<Integer>( 1 );
12     } // end BlockingBuffer constructor
13
14     // place value into buffer
15     public void set( int value ) throws InterruptedException
16     {
17         buffer.put( value ); // place value in buffer
18
19         System.out.printf( "%s%2d\t%s%d\n", "Producer writes ", value,
20             "Buffer cells occupied: ", buffer.size() );
21     } // end method set
22
23     // return value from buffer
24     public int get() throws InterruptedException
25     {
26         int readValue = buffer.take(); // remove value from buffer
27         System.out.printf( "%s %2d\t%s%d\n", "Consumer reads ",
28             readValue, "Buffer cells occupied: ", buffer.size() );
29         return readValue;
30     } // end method get
31 } // end class BlockingBuffer
```

Producer/Consumer Relationship: BlockingBufferTest

```
1  // Fig. 26.15: BlockingBufferTest.java
2  // Two threads manipulating a blocking buffer that properly
3  // implements the producer/consumer relationship.
4  import java.util.concurrent.ExecutorService;
5  import java.util.concurrent.Executors;
6
7  public class BlockingBufferTest
8  {
9      public static void main( String[] args )
10     {
11         // create new thread pool with two threads
12         ExecutorService application = Executors.newCachedThreadPool();
13
14         // create BlockingBuffer to store ints
15         Buffer sharedLocation = new BlockingBuffer();
16
17         application.execute( new Producer( sharedLocation ) );
18         application.execute( new Consumer( sharedLocation ) );
19
20         application.shutdown();
21     } // end main
22 } // end class BlockingBufferTest
```

Producer/Consumer Relationship: BlockingBufferTest

Producer writes	1	Buffer cells occupied:	1
Consumer reads	1	Buffer cells occupied:	0
Producer writes	2	Buffer cells occupied:	1
Consumer reads	2	Buffer cells occupied:	0
Producer writes	3	Buffer cells occupied:	1
Consumer reads	3	Buffer cells occupied:	0
Producer writes	4	Buffer cells occupied:	1
Consumer reads	4	Buffer cells occupied:	0
Producer writes	5	Buffer cells occupied:	1
Consumer reads	5	Buffer cells occupied:	0
Producer writes	6	Buffer cells occupied:	1
Consumer reads	6	Buffer cells occupied:	0
Producer writes	7	Buffer cells occupied:	1
Consumer reads	7	Buffer cells occupied:	0
Producer writes	8	Buffer cells occupied:	1
Consumer reads	8	Buffer cells occupied:	0
Producer writes	9	Buffer cells occupied:	1
Consumer reads	9	Buffer cells occupied:	0
Producer writes	10	Buffer cells occupied:	1

Producer done producing

Terminating Producer

Consumer reads	10	Buffer cells occupied:	0
----------------	----	------------------------	---

Consumer read values totaling 55

Terminating Consumer

Producer/Consumer Relationship with Synchronization

- Custom shared buffer with synchronization:

- Use of **synchronized** keyword and methods of class **Object**
- **Buffer** methods **get** and **set** are implemented as **synchronized methods**; a **thread** must obtain a **monitor lock** for the **buffer** to be able to operate on it
- Depending on whether **buffer** is in correct state or not, Object methods **wait**, **notify**, **notifyAll** are used to transition **threads** between **runnable** and **waiting** states
- Object method **wait** places calling **thread** in **waiting state** and releases **monitor lock** on the **buffer**, enabling any **thread** in **runnable state** to acquire the **lock**
- Object method **notify** allows a **thread** in **waiting state** to transition back to **runnable state**; method **notifyAll** allows all **threads** in **waiting state** to become **runnable**
- **IllegalMonitorStateException** results from a **thread** issuing **wait**, **notify**, **notifyAll** signals on an object without having a **monitor lock** to the object
- Using **notifyAll** ensures that all **waiting threads** become **runnable** eventually, to prevent thread **starvation**

Producer/Consumer Relationship Synchronized Buffer

```
1  // Fig. 26.16: SynchronizedBuffer.java
2  // Synchronizing access to shared data using Object
3  // methods wait and notifyAll.
4  public class SynchronizedBuffer implements Buffer
5  {
6      private int buffer = -1; // shared by producer and consumer threads
7      private boolean occupied = false; // whether the buffer is occupied
8
9      // place value into buffer
10     public synchronized void set( int value ) throws InterruptedException
11     {
12         // while there are no empty locations, place thread in waiting state
13         while ( occupied )
14         {
15             // output thread information and buffer information, then wait
16             System.out.println( "Producer tries to write." );
17             displayState( "Buffer full. Producer waits." );
18             wait();
19         } // end while
20
21         buffer = value; // set new buffer value
22
23         // indicate producer cannot store another value
24         // until consumer retrieves current buffer value
25         occupied = true;
26
27         displayState( "Producer writes " + buffer );
28
29         notifyAll(); // tell waiting thread(s) to enter runnable state
30     } // end method set; releases lock on SynchronizedBuffer
31 }
```

Producer/Consumer Relationship Synchronized Buffer

```

32 // return value from buffer
33 public synchronized int get() throws InterruptedException
34 {
35     // while no data to read, place thread in waiting state
36     while ( !occupied )
37     {
38         // output thread information and buffer information, then wait
39         System.out.println( "Consumer tries to read." );
40         displayState( "Buffer empty. Consumer waits." );
41         wait();
42     } // end while
43
44     // indicate that producer can store another value
45     // because consumer just retrieved buffer value
46     occupied = false;
47
48     displayState( "Consumer reads " + buffer );
49
50     notifyAll(); // tell waiting thread(s) to enter runnable state
51
52     return buffer;
53 } // end method get; releases lock on SynchronizedBuffer
54
55 // display current operation and buffer state
56 public void displayState( String operation )
57 {
58     System.out.printf( "%-40s%d\t\t%b\n\n", operation, buffer,
59         occupied );
60 } // end method displayState
61 } // end class SynchronizedBuffer

```

Producer/Consumer Relationship SharedBufferTest

```
1  // Fig. 26.17: SharedBufferTest2.java
2  // Two threads correctly manipulating a synchronized buffer.
3  import java.util.concurrent.ExecutorService;
4  import java.util.concurrent.Executors;
5
6  public class SharedBufferTest2
7  {
8      public static void main( String[] args )
9      {
10         // create a newCachedThreadPool
11         ExecutorService application = Executors.newCachedThreadPool();
12
13         // create SynchronizedBuffer to store ints
14         Buffer sharedLocation = new SynchronizedBuffer();
15
16         System.out.printf( "%-40s%s\t\t%s\n%-40s%s\n\n", "Operation",
17             "Buffer", "Occupied", "-----", "-----\t\t-----" );
18
19         // execute the Producer and Consumer tasks
20         application.execute( new Producer( sharedLocation ) );
21         application.execute( new Consumer( sharedLocation ) );
22
23         application.shutdown();
24     } // end main
25 } // end class SharedBufferTest2
```

Producer/Consumer Relationship SharedBufferTest

Operation -----	Buffer -----	Occupied -----
Consumer tries to read. Buffer empty. Consumer waits.	-1	false
Producer writes 1	1	true
Consumer reads 1	1	false
Consumer tries to read. Buffer empty. Consumer waits.	1	false
Producer writes 2	2	true
Consumer reads 2	2	false
Producer writes 3	3	true
Consumer reads 3	3	false
Producer writes 4	4	true
Producer tries to write. Buffer full. Producer waits.	4	true
Consumer reads 4	4	false
Producer writes 5	5	true
Consumer reads 5	5	false
Producer writes 6	6	true
Producer tries to write. Buffer full. Producer waits.	6	true
Consumer reads 6	6	false
Producer writes 7	7	true
Producer tries to write. Buffer full. Producer waits.	7	true
Consumer reads 7	7	false
Producer writes 8	8	true
Consumer reads 8	8	false
Consumer tries to read. Buffer empty. Consumer waits.	8	false

Producer/Consumer Relationship SharedBufferTest

Producer writes 9	9	true
Consumer reads 9	9	false
Consumer tries to read. Buffer empty. Consumer waits.	9	false
Producer writes 10	10	true
Consumer reads 10	10	false
Producer done producing Terminating Producer		
Consumer read values totaling 55 Terminating Consumer		

Producer/Consumer Relationship Bounded Buffers

- **Problem with producer/consumer program in section 26.7:**

We want the buffer to be of right size

- Application's performance significantly affected by relative speeds of Producer and Consumer threads
- If Producer thread produces much faster than Consumer thread can consume, Producer has to wait long until buffer has space to place next value
- When threads wait excessively, programs become less efficient, interactive programs become less responsive, applications suffer longer delays

- **Solution: bounded buffer:**

- Provides a fixed number of buffer cells into which the Producer thread can place values, and from which the Consumer thread can retrieve the values
- Minimizes the amount of waiting time for the threads sharing the buffer resource
- Buffer size should be optimized to minimize amount of thread wait time, while not wasting space

- **CircularBuffer:**

- Writes into and reads from the array elements in order, beginning at the first cell and moving towards the last
- When a Producer/Consumer reaches the last element, it returns to the first and begins writing/reading there

We basically want to reduce the program from consuming even though the producer is producing using a buffer will help to stop this

Producer/Consumer Relationship CircularBuffer

```
1 // Fig. 26.18: CircularBuffer.java
2 // Synchronizing access to a shared three-element bounded buffer.
3 public class CircularBuffer implements Buffer
4 {
5     private final int[] buffer = { -1, -1, -1 }; // shared buffer
6
7     private int occupiedCells = 0; // count number of buffers used
8     private int writeIndex = 0; // index of next element to write to
9     private int readIndex = 0; // index of next element to read
10
11     // place value into buffer
12     public synchronized void set( int value ) throws InterruptedException
13     {
14         // wait until buffer has space available, then write value;
15         // while no empty locations, place thread in blocked state
16         while ( occupiedCells == buffer.length )
17         {
18             System.out.printf( "Buffer is full. Producer waits.\n" );
19             wait(); // wait until a buffer cell is free
20         } // end while
21
22         buffer[ writeIndex ] = value; // set new buffer value
23
24         // update circular write index
25         writeIndex = ( writeIndex + 1 ) % buffer.length;
26
27         ++occupiedCells; // one more buffer cell is full
28         displayState( "Producer writes " + value );
29         notifyAll(); // notify threads waiting to read from buffer
30     } // end method set
31
32     // return value from buffer
33     public synchronized int get() throws InterruptedException
34     {
35         // wait until buffer has data, then read value;
36         // while no data to read, place thread in waiting state
37         while ( occupiedCells == 0 )
38         {
39             System.out.printf( "Buffer is empty. Consumer waits.\n" );
```

Producer/Consumer Relationship CircularBuffer

```
40         wait(); // wait until a buffer cell is filled
41     } // end while
42
43     int readValue = buffer[ readIndex ]; // read value from buffer
44
45     // update circular read index
46     readIndex = ( readIndex + 1 ) % buffer.length;
47
48     --occupiedCells; // one fewer buffer cells are occupied
49     displayState( "Consumer reads " + readValue );
50     notifyAll(); // notify threads waiting to write to buffer
51
52     return readValue;
53 } // end method get
54
55 // display current operation and buffer state
56 public void displayState( String operation )
57 {
58     // output operation and number of occupied buffer cells
59     System.out.printf( "%s%d\n", operation,
60         " (buffer cells occupied: ", occupiedCells, "buffer cells: " );
61
62     for ( int value : buffer )
63         System.out.printf( " %2d ", value ); // output values in buffer
64
65     System.out.print( "\n          " );
66
67     for ( int i = 0; i < buffer.length; i++ )
68         System.out.print( "---- " );
69
70     System.out.print( "\n          " );
71
72     for ( int i = 0; i < buffer.length; i++ )
73     {
74         if ( i == writeIndex && i == readIndex )
75             System.out.print( " WR" ); // both write and read index
76         else if ( i == writeIndex )
77             System.out.print( " W  " ); // just write index
78         else if ( i == readIndex )
79             System.out.print( " R  " ); // just read index
80         else
81             System.out.print( "    " ); // neither index
82     } // end for
83
84     System.out.println( "\n" );
85 } // end method displayState
86 } // end class CircularBuffer
```


Producer/Consumer Relationship CircularBufferTest

```
1 // Fig. 26.19: CircularBufferTest.java
2 // Producer and Consumer threads manipulating a circular buffer.
3 import java.util.concurrent.ExecutorService;
4 import java.util.concurrent.Executors;
5
6 public class CircularBufferTest
7 {
8     public static void main( String[] args )
9     {
10         // create new thread pool with two threads
11         ExecutorService application = Executors.newCachedThreadPool();
12
13         // create CircularBuffer to store ints
14         CircularBuffer sharedLocation = new CircularBuffer();
15
16         // display the initial state of the CircularBuffer
17         sharedLocation.displayState( "Initial State" );
18
19         // execute the Producer and Consumer tasks
20         application.execute( new Producer( sharedLocation ) );
21         application.execute( new Consumer( sharedLocation ) );
22
23         application.shutdown();
24     } // end main
25 } // end class CircularBufferTest
```

Producer/Consumer Relationship CircularBufferTest

```
Initial State (buffer cells occupied: 0)
buffer cells:  -1  -1  -1
               -----
               WR

Producer writes 1 (buffer cells occupied: 1)
buffer cells:   1  -1  -1
               -----
               R   W
```

```
Consumer reads 1 (buffer cells occupied: 0)
buffer cells:   1  -1  -1
               -----
               WR

Buffer is empty. Consumer waits.
Producer writes 2 (buffer cells occupied: 1)
buffer cells:   1   2  -1
               -----
               R   W

Consumer reads 2 (buffer cells occupied: 0)
buffer cells:   1   2  -1
               -----
               WR

Producer writes 3 (buffer cells occupied: 1)
buffer cells:   1   2   3
               -----
               W       R

Consumer reads 3 (buffer cells occupied: 0)
buffer cells:   1   2   3
               -----
               WR

Producer writes 4 (buffer cells occupied: 1)
buffer cells:   4   2   3
               -----
               R   W

Producer writes 5 (buffer cells occupied: 2)
buffer cells:   4   5   3
               -----
               R       W

Consumer reads 4 (buffer cells occupied: 1)
buffer cells:   4   5   3
               -----
               R   W

Producer writes 6 (buffer cells occupied: 2)
buffer cells:   4   5   6
               -----
               W       R

Producer writes 7 (buffer cells occupied: 3)
buffer cells:   7   5   6
               -----
               WR

Consumer reads 5 (buffer cells occupied: 2)
buffer cells:   7   5   6
               -----
               W       R

Producer writes 8 (buffer cells occupied: 3)
buffer cells:   7   8   6
               -----
               WR
```

Producer/Consumer Relationship Lock & Condition Interfaces

- **Lock and Condition interfaces:**

More precise control over thread synchronization

- Give more precise control over **thread synchronization**, but are *more complicated* to use
- Any object can contain a reference to an object that implements the **Lock interface** (of interface `java.util.concurrent.locks`)
- A **thread** calls **Lock** method **lock** to acquire the **lock**; only one **thread** can hold the **lock** at any time, others must wait
- A **thread** calls **Lock** method **unlock** to release the **lock**, allowing waiting **threads** to try acquire the **lock**
- Class **ReentrantLock**:
 - A basic implementation of the **Lock interface**
 - **fairness policy**: when **true**, “*longest-waiting thread acquires the lock when available*”; helps avoid *indefinite postponement*, but decreases performance
- **Condition object**:
 - Allows explicitly declaring the **condition** on which a **thread** will wait
 - Created by calling **Lock** method **newCondition**
 - More versatile than using *synchronized* keyword

Producer/Consumer Relationship Lock & Condition Interfaces

- **Lock and Condition interfaces:**
 - **Condition** methods **await**, **signal**, **signalAll** are called to **wait on** the **Condition** object and to let waiting **threads** transition to **runnable state**
 - **Deadlock**: occurs when a **waiting thread** (e.g. **thread1**) **cannot proceed** because it's **waiting for another thread** (e.g. **thread2**) to proceed, which also can't proceed because it's waiting on **thread1**
 - To **prevent starvation**, it's important to ensure that for every **await call** on a **Condition** object, there is a corresponding **signal call** to allow waiting **threads** to transition back to **runnable state**
 - **Lock** and **Condition** are powerful (allowing many things that can't be done using **Monitor lock**), but using them is also **error-prone** (e.g. not calling **unlock** can cause **starvation** for waiting **threads**)

Producer/Consumer Relationship Lock & Condition Interfaces

```
1 // Fig. 26.20: SynchronizedBuffer.java
2 // Synchronizing access to a shared integer using the Lock and Condition
3 // interfaces
4 import java.util.concurrent.locks.Lock;
5 import java.util.concurrent.locks.ReentrantLock;
6 import java.util.concurrent.locks.Condition;
7
8 public class SynchronizedBuffer implements Buffer
9 {
10     // Lock to control synchronization with this buffer
11     private final Lock accessLock = new ReentrantLock();
12
13     // conditions to control reading and writing
14     private final Condition canWrite = accessLock.newCondition();
15     private final Condition canRead = accessLock.newCondition();
16
17     private int buffer = -1; // shared by producer and consumer threads
18     private boolean occupied = false; // whether buffer is occupied
19
20     // place int value into buffer
21     public void set( int value ) throws InterruptedException
22     {
23         accessLock.lock(); // lock this object
24
25         // output thread information and buffer information, then wait
26         try
27         {
28             // while buffer is not empty, place thread in waiting state
29             while ( occupied )
30             {
31                 System.out.println( "Producer tries to write." );
32                 displayState( "Buffer full. Producer waits." );
33                 canWrite.await(); // wait until buffer is empty
34             } // end while
35
36             buffer = value; // set new buffer value
37
38             // indicate producer cannot store another value
39             // until consumer retrieves current buffer value
40             occupied = true;
41
42             displayState( "Producer writes " + buffer );
43
44             // signal any threads waiting to read from buffer
45             canRead.signalAll();
46         } // end try
47         finally
48         {
49             accessLock.unlock(); // unlock this object
50         } // end finally
51     } // end method set
52 }
```

Producer/Consumer Relationship Lock & Condition Interfaces

```
--
53 // return value from buffer
54 public int get() throws InterruptedException
55 {
56     int readValue = 0; // initialize value read from buffer
57     accessLock.lock(); // lock this object
58
59     // output thread information and buffer information, then wait
60     try
61     {
62         // if there is no data to read, place thread in waiting state
63         while ( !occupied )
64         {
65             System.out.println( "Consumer tries to read." );
66             displayState( "Buffer empty. Consumer waits." );
67             canRead.await(); // wait until buffer is full
68         } // end while
69
70         // indicate that producer can store another value
71         // because consumer just retrieved buffer value
72         occupied = false;
73
74         readValue = buffer; // retrieve value from buffer
75         displayState( "Consumer reads " + readValue );
76
77         // signal any threads waiting for buffer to be empty
78         canWrite.signalAll();
79     } // end try
80     finally
81     {
82         accessLock.unlock(); // unlock this object
83     } // end finally
84
85     return readValue;
86 } // end method get
87
88 // display current operation and buffer state
89 public void displayState( String operation )
90 {
91     System.out.printf( "%-40s%d\t\t%b\n\n", operation, buffer,
92         occupied );
93 } // end method displayState
94 } // end class SynchronizedBuffer
```

Producer/Consumer Relationship Lock & Condition Interfaces

```
1 // Fig. 26.21: SharedBufferTest2.java
2 // Two threads manipulating a synchronized buffer.
3 import java.util.concurrent.ExecutorService;
4 import java.util.concurrent.Executors;
5
6 public class SharedBufferTest2
7 {
8     public static void main( String[] args )
9     {
10         // create new thread pool with two threads
11         ExecutorService application = Executors.newCachedThreadPool();
12
13         // create SynchronizedBuffer to store ints
14         Buffer sharedLocation = new SynchronizedBuffer();
15
16         System.out.printf( "%-40s%s\t\t%s\n%-40s%s\n\n", "Operation",
17             "Buffer", "Occupied", "-----", "-----\t\t-----" );
18
19         // execute the Producer and Consumer tasks
20         application.execute( new Producer( sharedLocation ) );
21         application.execute( new Consumer( sharedLocation ) );
22
23         application.shutdown();
24     } // end main
25 } // end class SharedBufferTest2
```

Producer/Consumer Relationship Lock & Condition Interfaces

Operation -----	Buffer -----	Occupied -----
Producer writes 1	1	true
Producer tries to write. Buffer full. Producer waits.	1	true
Consumer reads 1	1	false
Producer writes 2	2	true
Producer tries to write. Buffer full. Producer waits.	2	true
Consumer reads 2	2	false
Producer writes 3	3	true
Consumer reads 3	3	false
Producer writes 4	4	true
Consumer reads 4	4	false
Consumer tries to read. Buffer empty. Consumer waits.	4	false
Producer writes 5	5	true
Consumer reads 5	5	false
Consumer tries to read. Buffer empty. Consumer waits.	5	false
Producer writes 6	6	true
Consumer reads 6	6	false
Producer writes 7	7	true
Consumer reads 7	7	false
Producer writes 8	8	true
Consumer reads 8	8	false
Producer writes 9	9	true
Consumer reads 9	9	false
Producer writes 10	10	true
Producer done producing Terminating Producer		
Consumer reads 10	10	false
Consumer read values totaling 55 Terminating Consumer		